



simulate once



$m \sim p(m)$;
fine-grid trajectories



train once per system



≈ 1 h, one GPU



amortized inference

any observation set:
any K , any placement

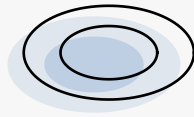


posterior draws, ≈ 70 ms



evaluation

validated reference
on the same points



FID