

Amortized Bayesian parameter recovery for dynamical systems with conditional flow matching

The `amortix` package: architecture, evaluation methodology, and benchmarks

Ivan Oseledets

August 2026 (v0.1)

Abstract

We describe `amortix`, a package for amortized Bayesian estimation of the parameters of dynamical systems — ordinary and stochastic differential equations and discretized PDEs — from noisy observations at a finite set of time points. A conditional flow-matching model is trained once on simulated (parameter, data) pairs and thereafter produces posterior samples for any new dataset without likelihood evaluations or per-dataset optimization; a single trained network handles any number and placement of observation points, the setting required for sequential experiment refinement and Bayesian experimental design. The package couples three components that we argue are only useful together: (i) a set-transformer encoder whose input representation is learnable rather than prescribed — a monotone mixture-of-scales reparametrization of the value axis, attention phases computed from physical observation times, and explicit set-cardinality features; (ii) a benchmark suite of fifteen systems spanning finance, neuroscience, epidemiology, ecology, pharmacokinetics, classical mechanics, and reaction-diffusion PDEs; and (iii) an evaluation harness in which simulation-based calibration is used only as a screen, and the instrument of record is comparison against exact or validated Markov chain Monte Carlo reference posteriors on the same observed points. On the fixed-design suite, 31 of 32 parameters across nine systems pass simulation-based calibration with a mean coverage error of 1.3 percentage points; on the variable-design suite, posterior widths match exact references to within 5–7% uniformly over design sizes $K = 5\text{--}100$, and the two residual defects found anywhere in the suite decay as measured power laws of the training budget. We compare against neural posterior estimation and flow-matching posterior estimation from the `sbi` package, both network families of BayesFlow 2, and a port of the Simformer, on three systems with exact or validated references. On a multiplicative process, all external configurations either inflate the volatility posterior by factors of 1.4–3.2 or become overconfident (widths 0.4–0.9) depending on the input coordinates they are given, while the learnable representation recovers the exact posterior from raw data. On an additive control process the same external tools improve markedly — the systematic bias disappears and the width inflation shrinks from 2.5–3.2 to 1.6–1.9 — which localizes the dominant failure to input conditioning while showing that a residual representation gap remains even in the benign case. We document this failure mechanism — the conditioning of the encoder input — in three distinct forms, together with learnable modules that repair each one.

1 Introduction

Recovering the parameters of a differential-equation model from noisy measurements of its solution is one of the basic inverse problems of applied science: a modeler writes down a dynamical system — an ODE, an SDE, or a PDE — whose right-hand side depends on unknown physical constants, observes a trajectory at a finite set of time points, and asks what the observations say about the constants. Like most inverse problems, this one is ill-posed in the practical sense: different parameter values can produce nearly indistinguishable observations, especially when the state is partially observed or the sampling is sparse, so a point estimate both hides the ambiguity and is unstable under noise. The Bayesian formulation resolves this by reporting the full posterior distribution over parameters given the data — ambiguity becomes posterior spread and correlation, quantities one can act on — at the price of two computational obstacles. First, for most stochastic and partially observed systems, the likelihood of the observations requires integrating over unobserved trajectory segments and is intractable or expensive. Second, modern workflows need posteriors *repeatedly*: calibration studies process thousands

of synthetic datasets, sequential experiments update the posterior after every new measurement, and Bayesian experimental design compares candidate observation schedules by the posteriors they would produce — so per-dataset MCMC or optimization inside such loops is the dominant cost even when a likelihood is available.

Simulation-based inference (SBI) addresses both issues by learning a conditional generative model of the parameters given the data from simulated pairs — parameters drawn from the prior, datasets produced by the simulator — thereby amortizing the cost of inference into a one-time training phase (Cranmer et al., 2020). Neural posterior estimation (NPE) and its sequential variants (Papamakarios and Murray, 2016; Greenberg et al., 2019) fit normalizing flows by maximum likelihood; flow-matching and diffusion variants (Wildberger et al., 2023; Gloeckler et al., 2024) replace the density estimator with transport-based generative models. Mature software exists: the **sbi** toolkit (Boelts et al., 2024), BayesFlow (Radev et al., 2020), and the **sbibm** benchmark collection (Lueckmann et al., 2021).

This report describes **amortix**, a compact PyTorch package (about 3,500 lines; one default model of 4.8×10^5 weights) for amortized Bayesian parameter recovery in dynamical systems. Three aspects of its design distinguish it from the toolkits above, and the report is organized around measuring their consequences.

The first concerns the interface between raw data and the network. Conditional density estimators are conventionally fitted to a fixed-length data vector after global standardization; the estimator sees whatever coordinates the user chose. We show in Sec. 8.5 that on multiplicative and heavy-tailed processes this convention fails in a specific, reproducible way across four different generative engines from three codebases, and that supplying hand-derived coordinates does not repair it uniformly: with identical inputs, different implementations end up over-dispersed, overconfident, or biased. In **amortix** the transformation from observations to network inputs is itself part of the trained model — a monotone mixture-of-scales reparametrization of the value axis, observation times entering attention as continuous phases, and explicit design-size features (Sec. 4); each component is motivated by a measured failure and validated by an ablation.

The second concerns observation designs. **amortix** trains a single network to approximate $p(m | S)$ coherently for every subset S of the measurement points of an experiment — arbitrary sizes and placements, including spatio-temporal sensor designs for PDEs (Sec. 8.3). This is the object consumed by sequential refinement and Bayesian experimental design, and supporting it changes the training procedure (Sec. 5): designs must be resampled fresh at every optimizer step, and the distribution of design sizes becomes a modeling choice with measurable consequences.

The third concerns evaluation. Simulation-based calibration (Talts et al., 2018) is a necessary-condition test whose statistical power at practical sample sizes is limited; we measure that limit (Sec. 6) and use SBC only as a screen. The instrument of record throughout is the comparison of posterior mean and width against an exact or validated-MCMC reference posterior on the same observed points, and the package ships this machinery — closed forms, exact-likelihood adaptive Metropolis samplers, chain validation — as part of its public API. The concern motivating this discipline is documented in the SBI literature: learned posteriors can be confidently wrong in ways that rank diagnostics do not detect (Hermans et al., 2022).

The remainder of the report presents the method (Sec. 3), the architecture (Sec. 4), training procedures (Sec. 5), the evaluation methodology (Sec. 6), the benchmark suite (Sec. 7), experimental results including comparisons with **sbi**, BayesFlow, and the Simformer on three systems with exact or validated references (Sec. 8), an analysis of the failure modes encountered during development (Sec. 9), and limitations (Sec. 10). Appendix A defines every benchmark system with its equations, priors, and observation model; Appendix B details the reference posteriors; Appendix C maps the terminology of this report to package options.

2 Problem formulation

Let a dynamical system depend on a parameter vector $m \in \mathbb{R}^d$ with a known prior $p(m)$, uniform on a box in all benchmarks below. The system produces a state trajectory $x(t)$ — through an ODE $\dot{x} = f(x; m)$, an SDE $dx = f(x; m)dt + \sigma(x; m)dW_t$, or a spatially discretized PDE — and the trajectory is observed through a noisy channel at K points:

$$y_i = h(x(t_i)) + \varepsilon_i, \quad \mathcal{D} = \{(t_i, y_i)\}_{i=1}^K. \quad (1)$$

The observation operator h may expose only part of the state: in the benchmarks below, only the price of an asset while its stochastic volatility is latent, only the membrane voltage of a neuron while its

gating variables are latent, or only three spatial sensors of a PDE field. The target is the posterior $p(m \mid \mathcal{D}) \propto p(\mathcal{D} \mid m)p(m)$.

Three standing assumptions delimit the setting. The parameter is finite- and low-dimensional ($d = 2\text{--}5$ in every benchmark below) with a prior that factorizes across components and has tractable marginal CDFs (uniform boxes throughout; the construction extends to any such prior). The simulator can be run forward cheaply and repeatedly, but nothing else is required of it — no likelihood evaluations, no gradients. Observation times (and sensor indices) are known exactly; the measurement noise ε_i is independent across observations with a known per-system law.

Running example. The Ornstein–Uhlenbeck process is used throughout to make constructions concrete: $dX = -\theta X dt + \sigma dW$ with $m = (\theta, \sigma)$, prior box $[0.3, 3] \times [0.2, 1.5]$, stationary start $X_0 \sim \mathcal{N}(0, \sigma^2/2\theta)$. A dataset is, for instance, $K = 74$ noisy values of X on a fixed two-rate grid, or — in the variable-design setting — any K values at arbitrary times. Its posterior has the features that recur across the suite: the diffusion coefficient σ is identified sharply by high-frequency increments; the reversion rate θ only weakly over a finite horizon; and the width of both marginals must shrink as observations accumulate. An estimator can therefore be wrong in the posterior mean, in the posterior width, or in the width-versus- K trend, and the evaluation methodology of Sec. 6 is built to detect all three.

Although the posterior itself is intractable, sampling from the *joint* distribution of parameters and data is easy: draw $m_j \sim p(m)$, simulate the system, apply the observation channel. This yields arbitrarily many training pairs $(m_j, \mathcal{D}_j)$, and a conditional generative model trained on them approximates the posterior for every dataset simultaneously — the cost of inference is moved into a one-time training phase, and inference for a new dataset is a forward pass. We impose one requirement beyond standard amortization: a single trained network must remain correct for any number K and any placement of observation times (and, for PDEs, sensor positions). Formally, for every subset S of the measurement points of one experiment the network must approximate $p(m \mid S)$, and the family of answers must be coherent — in particular the posterior must widen as points are removed. We refer to this requirement as *design amortization*; it is the object consumed by sequential refinement and by Bayesian experimental design, where candidate observation schedules are ranked by the posteriors they would produce.

3 Posterior sampling by conditional flow matching

The idea of the method is to represent the posterior not by a density formula but by a *transport*: a learned, data-conditioned deformation that carries a fixed reference distribution onto the posterior. Sampling then amounts to drawing from the reference and applying the deformation — there is no normalizing constant to compute and no invertibility constraint to build into the network, because the deformation is defined implicitly as the flow of a velocity field, and the field itself is fitted by ordinary least-squares regression on simulated pairs. This construction is conditional flow matching (Lipman et al., 2023).

Concretely, parameters are first mapped componentwise to normalized coordinates

$$z = \Phi^{-1}(F_{\text{prior}}(m)) \in \mathbb{R}^d, \quad (2)$$

where F_{prior} collects the marginal prior CDFs and Φ is the standard normal CDF; in z -coordinates the prior is exactly $\mathcal{N}(0, I)$, box constraints disappear, and — usefully — the uninformative limit costs nothing: when the data carry no information the posterior coincides with the reference and the transport is the identity. Let $z_0 \sim \mathcal{N}(0, I)$ be a reference draw and z_1 the normalized parameter drawn jointly with its dataset $\mathcal{D}$. Define the linear interpolation $z_t = (1 - t)z_0 + tz_1$ and train a velocity network by regression:

$$\mathcal{L}(\theta) = \mathbb{E}_{m, \mathcal{D}, z_0 \sim \mathcal{N}(0, I), t \sim U[0, 1]} \|v_\theta(z_t, t, \mathcal{D}) - (z_1 - z_0)\|^2. \quad (3)$$

The minimizer of (3) is the marginal velocity field whose flow transports $\mathcal{N}(0, I)$ at $t = 0$ to $p(z_1 \mid \mathcal{D})$ at $t = 1$ (Lipman et al., 2023). Sampling integrates $\dot{z} = v_\theta(z, t, \mathcal{D})$ with a fixed-step midpoint scheme and maps back through the inverse of (2); the number of steps is a solver-accuracy parameter fixed once by a convergence check — refining the default (20 steps) to fourth-order Runge–Kutta with three times as many steps changes no benchmark statistic (Sec. 8.4). Evaluation uses an exponential moving average of the weights. The same estimator family, with a different conditioning architecture, is available in `sbi` as FMPE (Wildberger et al., 2023); the comparison of Sec. 8.5 therefore isolates the effect of the input representation from the choice of generative engine.

No likelihood evaluations, simulator gradients, or per-dataset optimization are required at any stage; the simulator is executed only forward, during training-set generation.




Figure 1: Model architecture; arrows denote data flow, bottom to top. Left: each observation is featurized and linearly embedded into one token; L transformer encoder blocks, whose attention phases are computed from the physical observation times (Sec. 4.2), produce a memory H of K encoded states. Middle: the d components of the flow state z_t form d parameter-slot tokens processed by L velocity blocks and a linear head into the velocity v_θ . Right: one velocity block expanded — self-attention across parameter slots, cross-attention reading the observation memory H , and an MLP, each preceded by layer normalization with scale-shift (γ, β) and followed by a gated residual connection (g , marked $+$); the triples (γ, β, g) are produced from the flow time t by a zero-initialized map, so the untrained block is the identity (Peebles and Xie, 2023).

The training objective is consistent in the following sense.

Proposition 1 (consistency). *Let $(m, \mathcal{D}) \sim p(m)p(\mathcal{D} \mid m)$ and let z_1 be given by (2). If the loss (3) is minimized over all measurable velocity fields, then for $p(\mathcal{D})$ -almost every $\mathcal{D}$ the flow of the minimizer transports $\mathcal{N}(0, I)$ at $t = 0$ to the conditional law $p(z_1 \mid \mathcal{D})$ at $t = 1$; equivalently, the sampling procedure above draws exactly from the posterior (Lipman et al., 2023; Wildberger et al., 2023).*

The proposition is an idealization, and the two ways in which practice departs from it shape the rest of the paper. With a finite network, a finite simulation budget, and a numerical solver, no calibration guarantee survives; the residual error is an empirical quantity, and Sec. 6 builds the instruments to measure it rather than assume it away. A second, structural observation: the loss (3) is a plain average over simulated data and flow times, and nothing in it privileges accuracy in directions of parameter space where the posterior is nearly flat — errors there change the regression target little. We offer this as an interpretation, not a theorem, and it is consistent with what the experiments show: residual posterior widths sit slightly above the reference and decrease with training budget, and the largest residual bias we observe anywhere in the suite lies along the flat ridge direction of the Fisher–KPP posterior (Sec. 8.4).

4 Conditioning architecture

The velocity field $v_\theta(z, t, \mathcal{D})$ must be conditioned on a dataset that is a *set of time-stamped measurements*, and three properties of that set drive the architecture: the number of observations K varies from dataset to dataset (by two orders of magnitude in the variable-design setting); the observation times are irregular real numbers, not grid indices; and the times carry all ordering information, so the output must not depend on how the observations happen to be stored. These properties rule out the fixed-length-vector conditioners that are the default in NPE implementations, and they are naturally met by a transformer operating on one feature vector per observation — throughout this section, a *token*.

4.1 Overall scheme

The network evaluates $v_\theta(z, t, \mathcal{D})$, where z is the normalized parameter vector partway along its path from reference noise to posterior sample — the interpolation z_t at training time, the current ODE state at sampling time. The design follows the encoder–decoder pattern that is standard in conditional flow-matching and diffusion models: the data are encoded once, and a decoder combines that encoding with the object being generated,

$$v_\theta(z, t, \mathcal{D}) = \text{Dec}(z, t, \text{Enc}(\mathcal{D})). \quad (4)$$

Two differences from the image and text instances of this pattern dictate the concrete form (Fig. 1). First, the conditioning is not a fixed-size embedding but a set whose size K varies: Enc is therefore a transformer over per-observation tokens, its output is a *memory* of K encoded states rather than one pooled vector, and Dec reads the memory by cross-attention — the one standard mechanism indifferent to K . Second, the generated object is not an image but a short vector of d parameters, so the decoder is small: the d components of z are treated as d tokens, self-attention among them expresses posterior correlations (the ridge of the Fisher–KPP posterior, the drift–volatility coupling of GBM), and the flow time t enters every block through zero-initialized adaptive layer normalization (adaLN-Zero), the conditioning mechanism of diffusion transformers (Peebles and Xie, 2023), so that the untrained network is an identity-like map.

On the encoder side, observation i is described by a few features — its time t_i , its transformed value $w(y_i)$ with w the learnable monotone map of Sec. 4.3, an index identifying the sensor or state component that produced the value (multivariate states contribute one token per observed component per time), and the design-size feature of Sec. 4.3 — and lifted to the model width by a learned linear map. The attention mechanism that lets these tokens interact is the subject of the next subsection; the choice of features is the subject of Sec. 4.3.

This scheme was not the first attempt but the survivor of a sequence of measured alternatives; the failed ones are retained in the package as reproducible controls and are compared under a common protocol in the ablation study of Sec. 8.6. All sizes are configuration parameters; the package default, used unchanged for every experiment in this paper, is model width 64, four attention heads, and three blocks on each side (4.8×10^5 trainable parameters). The default is deliberately small: Sec. 8.4 shows that doubling width and depth changes no benchmark result, so at these problem sizes accuracy is limited by training budget, not capacity.

4.2 Attention phases from observation times

Attention compares tokens through inner products: the logit between tokens i and l in one head is $q_i^\top k_l / \sqrt{d_h}$, where the query q_i and key k_l are linear functions of the tokens and d_h is the head dimension. Rotary position embeddings (Su et al., 2024) make these logits position-aware as follows: the d_h coordinates of each query and key are grouped into $d_h/2$ two-dimensional planes, and in plane j both vectors are rotated by

$$R(\phi) = \begin{pmatrix} \cos \phi & -\sin \phi \\ \sin \phi & \cos \phi \end{pmatrix}, \quad \phi = \omega_j n, \quad (5)$$

where n is the token’s integer position and the frequencies ω_j decrease geometrically across planes (fast planes resolve near neighbors, slow planes long ranges). Since rotations preserve inner products up to the relative angle,

$$(R(\omega_j n) q)^\top R(\omega_j n') k = q^\top R(\omega_j (n' - n)) k, \quad (6)$$

attention logits depend only on position differences $n' - n$.

For irregularly sampled data the storage position n is meaningless, and the natural continuous substitute is the physical observation time — the same substitution that attention models for irregular clinical and event time series make in other positional schemes (Horn et al., 2020; Shukla and Marlin, 2021). We therefore set $\phi = \omega_j t_i$, upon which (6) makes every attention logit a function of the time difference $t_i - t_l$ and of nothing else about storage order. The encoder is then permutation-invariant by construction (verified to 10^{-8} under random token permutations), and design amortization involves no positional bookkeeping.

The measurement that forced this choice is worth recording. An earlier configuration used integer-position rotations, trained with long token sequences subsampled in place (kept tokens retained their spread-out positions), and evaluated on contiguously packed short sequences. Every component was individually correct, yet the mismatch between the positional distributions seen at training and evaluation produced systematic posterior biases of +1.1 to +3.3 standard deviations (measured in Sec. 8.6). Integer

positions create an implicit train/test contract on token placement; time-based phases eliminate the contract rather than manage it.

4.3 Input features

The remaining design choices concern what the tokens contain. Three of them matter and are described in turn: a learnable transform of the observed values, an optional increment-based token construction for directly observed diffusions, and an explicit design-size feature. Each occupies a few lines of code; each is load-bearing, as the ablations of Sec. 8.6 show.

Value transform. The values y_i enter tokens through a monotone map

$$w(x) = \sum_{k=1}^8 a_k \operatorname{sign}(x) \log(1 + |x|/s_k), \quad a_k \geq 0, \quad (7)$$

a positive mixture over a geometric bank of S scales s_k ; the same construction (with its own weights) is applied to time gaps where they appear. Since $\log(1 + |x|/s) \approx |x|/s$ for $|x| \ll s$ and $\approx \log |x| - \log s$ for $|x| \gg s$, each component is linear below its scale and logarithmic above it, so the mixture can realize the identity, the signed logarithm, and all intermediate compressions, with monotonicity guaranteed by the sign constraint on the weights. The number of scales S and their placement are hyperparameters with a transparent role — the bank must cover the dynamic range of the observed values — and we use one default everywhere, $S = 8$ scales spaced by octaves, for all fifteen benchmark systems with no per-problem adjustment. On geometric Brownian motion the trained mixture concentrates on its small-scale components: it learns the log-price coordinate in which increments become homoscedastic.

The structure of (7) — monotone by construction, few parameters, containing the useful coordinates at nonzero gradient from initialization — is the point of the design; less structured learnable alternatives were evaluated and did not work at matched budget. A trainable Yeo–Johnson transform (Yeo and Johnson, 2000), which nests the logarithm at $\lambda \rightarrow 0$, never left its initialization on GBM (the volatility bias stayed at +0.48 posterior sd, as with no transform at all), and unconstrained feature maps — a two-layer MLP on raw values, and random Fourier features — also failed to recover the log coordinate.

Increment features for directly observed diffusions. For scalar SDEs whose state is observed directly — no latent components — the package offers a second token construction: after sorting by time, consecutive observations contribute

$$[t_{(i)}, w(y_{(i)}), \Delta w_i, (\Delta w_i)^2, w_t(\Delta t_{(i)})], \quad \Delta w_i = w(y_{(i+1)}) - w(y_{(i)}).$$

The motivation is elementary: for a Markov observed process the likelihood factorizes over consecutive pairs, transition densities of diffusions depend on increments and gaps, and squared increments carry the quadratic-variation information that identifies diffusion coefficients. This is an inductive bias of the same character as convolutions for images — useful where its assumption holds, not a requirement of the method: the point-token path of Sec. 4.1 is fully general, and the ablations of Sec. 8.6 quantify what the increment tokens buy at fixed budget.

The rule for choosing between the constructions — increment tokens exactly when the observed process is Markov — was tested in both directions: an increment-token model class transfers across Markov systems without adjustment (GBM to OU), and on the Heston model, whose observed price is not Markov because volatility is latent, increment and point tokens perform identically across all 25 evaluation cases, as the factorization argument predicts. The package selects the construction from a per-problem flag.

Two additions make increment tokens robust on heavy-tailed processes (the measurements are in Sec. 9). Squared increments are logarithmically compressed before feature normalization. And the normalization is made dataset-adaptive by feature-wise affine modulation (FiLM; Perez et al., 2018): pooled robust statistics s of the token set (computed over non-padding positions only) are mapped by a zero-initialized MLP to per-channel scales and shifts,

$$h_c \mapsto (1 + \gamma_c(s)) h_c + \beta_c(s), \quad (8)$$

applied after the running feature normalization. Zero initialization keeps the working baseline at the start of training; the ordering (normalize, then modulate) matters and is discussed with the failure it prevents in Sec. 9.

Design size. Posterior width is governed by the number of observations — in regular models the error scales as $K^{-1/2}$ — but K is a global property of the set, and softmax attention, which computes normalized weighted means, is nearly invariant to it. Softmax attention computes normalized weighted means, which are nearly invariant to the number of tokens — yet posterior width is governed by K . The deficit is measurable by linear probes from the pooled encoder output: all architecture variants of Table 8 recover posterior-mean statistics ($R^2 = 0.83\text{--}0.94$), while $R^2(\log K)$ ranges from 0.985 to 0.498 and ranks the variants exactly in the order of their width errors. The repair is to make the count available, and the package default is the simplest form: $\log K$ enters every token as one of the input features. An alternative that requires no explicit feature — appending a few learned auxiliary tokens to the set, in the spirit of register tokens (Darcet et al., 2024), and letting the model accumulate the count in them — achieves the same effect in our ablations; the number of such tokens is a hyperparameter with no measurable influence beyond one. With the count available, the width error on the hardest case ($K = 100$) drops from 1.95 to 1.49 with bias unchanged; the residual is closed by the training procedure of the next section.

5 Training procedures

Common settings. All models are trained with Adam at a base learning rate of 3×10^{-4} under a cosine schedule with linear warmup, batch size 64, and gradient clipping; evaluation uses an exponential moving average of the weights with decay 0.999. At every optimizer step the flow times t are drawn afresh (stratified over $[0, 1]$) together with fresh reference noise z_0 , so no pairing between datasets and flow times is ever repeated. For fixed-design problems the training set consists of n simulated datasets and training runs for a fixed number of epochs over it; the gallery of Sec. 8.1 uses $n = 40,000$ simulations and 35 epochs.

Design amortization. Training data are simulated once on a fine grid; observation designs are then resampled *from the stored fine trajectories at every optimizer step*. The procedure has three components, each fixed by a controlled comparison on three systems (GBM, Ornstein–Uhlenbeck, FitzHugh–Nagumo):

1. *Fresh designs at every step.* Re-drawing each trajectory’s observation subset at each step makes design memorization impossible and lets one simulation serve arbitrarily many designs. The alternative — one fixed design per trajectory — closes the dense-design error at long budgets but overfits the design set, producing overconfidence at sparse K (width ratio 0.86 against the exact reference).
2. *The design-size law.* K is drawn 50% log-uniform on $[K_{\min}, K_{\max}]$ and 50% uniform on the dense half of the range. Log-uniform sampling alone places $\approx 6\%$ of the mass at $K \geq 100$ and starves the dense regime, which appears as inflated widths precisely there; the mixture alone halves the dense-design excess (GBM $1.27 \rightarrow 1.14$, OU $1.34 \rightarrow 1.14$; the FitzHugh–Nagumo dense- K SBC failure at $p < 10^{-3}$ disappears, $p = 0.54$).
3. *Budget scaling.* When higher accuracy is required, simulations and optimizer steps are scaled together at constant epochs; Sec. 8.4 shows the residuals then follow power laws in this single budget variable.

With this procedure, the GBM posterior width matches the exact per-design posterior to 5–7% uniformly over $K = 5, 9, 20, 100$ (Fig. 7).

6 Evaluation methodology

Two instruments with different roles are used throughout, plus a diagnostic probe.

6.1 Simulation-based calibration as a screen

SBC (Talts et al., 2018) tests a necessary condition: for $m^* \sim p(m)$ and $\mathcal{D} \sim p(\mathcal{D} \mid m^*)$, the rank of each component of m^* among posterior samples must be uniform. We run SBC per parameter (uniformity test on 500 simulated datasets with 200 posterior draws each unless stated; $p > 0.05$ passes) and report coverage of central 50% and 90% intervals.

The power of the test at this size is limited, and because our conclusions depend on knowing that limit, we measured it: on cases where an exact reference is available, SBC at 300–500 datasets fails to reject posterior width ratios up to ≈ 1.3 , and its per-parameter p -values are unstable between retrainings

when residual biases are in the range 0.1–0.2 posterior standard deviations. SBC also cannot penalize a posterior that ignores the data and returns the prior; the prior-mean and ridge controls of Sec. 7 exist for that reason. SBC is therefore used as a cheap screen, never as a verdict.

6.2 Reference posteriors as the record

Wherever a tractable likelihood exists, we compare distributions on the same observed points using two statistics per parameter: the signed error of the posterior mean in units of the reference posterior standard deviation (*bias*) and the ratio of standard deviations (*width*; 1 is exact, > 1 conservative, < 1 overconfident). References are of three kinds (full details in Appendix B):

- *Closed form.* The linear-Gaussian benchmark (exact Gaussian posterior) and geometric Brownian motion, whose log-returns give a conjugate normal-inverse- χ^2 posterior importance-corrected to the box prior.
- *Exact-likelihood MCMC.* An adaptive random-walk Metropolis sampler (after Haario et al., 2001): proposal scale adapted to a 0.25 acceptance target during burn-in, proposal shape replaced by the empirical covariance Cholesky factor rescaled to $2.38/\sqrt{d}$, adaptation frozen before the retained phase, out-of-box proposals rejected (which samples the box-truncated target exactly). Likelihoods are exact per-gap transition densities: Gaussian Euler–Maruyama chain transitions for OU (matching the simulator that generated the data, plus the stationary density of the informative initial state), a Poisson-mixture transition density for the Merton jump-diffusion, log-normal residuals around the Bateman curve for pharmacokinetics, and a Gaussian observation model around the deterministic PDE solve for Fisher–KPP — in the last case one likelihood evaluation is one PDE solve.
- *Chain validation.* A reference is used only after validation. Every MCMC reference in this report is run at least twice from well-separated starting points, and the between-chain discrepancy of posterior means (in posterior-sd units) is reported with the result it certifies. Where a closed form and a chain both exist, the chains reproduce it to a worst mean discrepancy of 0.099 sd and width ratios 0.91–1.03 over 24 comparisons.

Twice during development the screen and the record contradicted each other; both times the contradiction exposed an implementation error (a padding mask leaking into set-level statistics; the positional contract of Sec. 4.2). We consider the redundancy a feature of the methodology rather than an inefficiency.

6.3 Encoder probing

To separate encoder failures from flow failures, we fit ridge regressions from the pooled encoder output to known sufficient statistics of the problem. High probe R^2 with poor posterior accuracy localizes the defect downstream of extraction. In the Merton analysis of Sec. 9, the failing configuration extracted robust volatility statistics *better* than the healthy one ($R^2 = 0.967$ on truncated realized variance), redirecting the repair from the encoder’s capacity to the stability of its input representation.

7 Benchmark suite

The suite contains fifteen systems (Table 1): nine with a fixed observation design used for the calibration gallery, and six with variable designs used for design amortization. Every system ships with a classical estimator (`sota()`) and, where a tractable likelihood exists, a reference posterior. Full equations, priors, and observation models are given in Appendix A.

Controls and classical baselines. Every result is read against: (i) a prior-mean predictor (exactly 25% mean absolute error under a uniform prior) and a degree-2 ridge regression on ~ 20 generic summary statistics — these two detect prior-returning posteriors that SBC passes; (ii) the per-system classical estimator: exact MLE for OU and GBM, Euler pseudo-likelihood for CIR, Kramers–Moyal and SINDy-style regression for drift discovery (Brunton et al., 2016), multi-start nonlinear least squares for FitzHugh–Nagumo, SEIR, and Lotka–Volterra, and the exact posterior mean for the linear-Gaussian case; (iii) the reference posteriors of Sec. 6.2.

system	type	d	obs.	design	reference posterior
linear-Gaussian	linear map	4	full	fixed	exact Gaussian
Ornstein–Uhlenbeck	SDE	2	full	fixed	exact-likelihood MCMC
geometric Brownian motion	SDE	2	full	fixed	conjugate closed form
Cox–Ingersoll–Ross	SDE	3	full	fixed	— (Euler pseudo-MLE)
double well	SDE	3	full	fixed	—
stochastic Lotka–Volterra	2-d SDE	4	full	fixed	—
FitzHugh–Nagumo	2-d SDE	4	v only	fixed	—
SEIR	4-d SDE	5	I, D	fixed	—
SINDy-SDE (cubic drift)	SDE	5	full	fixed	—
Heston	2-d SDE	5	price only	variable	—
Merton jump-diffusion	jump SDE	5	full	variable	Poisson-mixture MCMC
Hénon–Heiles	Hamiltonian ODE	3	q_1 only	variable	—
Hodgkin–Huxley	4-d ODE	4	V only	variable	—
pharmacokinetics (Bateman)	ODE (closed form)	3	concentration	variable	log-normal MCMC
Fisher–KPP	PDE (64-point grid)	2	3 sensors	variable	PDE-likelihood MCMC

Table 1: The benchmark suite. “obs.” lists the observed coordinate when the state is only partially visible; “design” distinguishes the fixed-grid gallery from variable-design systems, where every dataset carries its own set of (time, sensor) observation points. Dashes in the last column indicate systems where no tractable reference exists — these are evaluated by SBC and controls only, which is precisely the regime that motivates carrying exact-reference systems in the same suite.

system	evaluation	amortix	externals	details
linear-Gaussian	exact posterior	SBC 4/4; width ≈ 1.00	—	Tab. 3
Ornstein–Uhlenbeck	exact-lik. MCMC	1.03 / 1.05; SBC 2/2	NPE 1.88; FMPE 1.60	Tab. 6
geometric BM	conjugate closed form	1.00 / 1.05; SBC 2/2	1.44–3.16 (11 arms)	Tab. 4, 5
Cox–Ingersoll–Ross	SBC + controls	2/3 (see text)	—	Tab. 3
double well	SBC + controls	3/3	—	Tab. 3
stoch. Lotka–Volterra	SBC + controls	4/4	—	Tab. 3
FitzHugh–Nagumo	SBC + controls	4/4	—	Tab. 3
SEIR	SBC + controls	5/5	—	Tab. 3
SINDy-SDE	SBC + controls	5/5	—	Tab. 3
Heston	SBC (variable designs)	all 25 cases pass	—	Sec. 8.3
Merton jump-diff.	Poisson-mixture MCMC	widths $\rightarrow 1.06$ w/ budget	—	Fig. 5, 3
Hénon–Heiles	SBC (variable designs)	pass at all $K \geq 6$	—	Sec. 8.3
Hodgkin–Huxley	SBC (variable designs)	4/4 mixed designs	—	Sec. 8.3
pharmacokinetics	log-normal MCMC	$ \text{bias} \leq 0.06$; width 1.09–1.33	—	Fig. 3
Fisher–KPP (PDE)	PDE-likelihood MCMC	1.02 / 1.04 (fixed design)	NPE 1.28; FMPE fails D	Tab. 7

Table 2: Overview of all quantitative results. The “externals” column quotes the width ratio of the least well-conditioned parameter for each external method (GBM row: the range across the eleven external arms of Tables 4 and 5). SBC counts are parameters passing the uniformity screen at $p > 0.05$. Dashes: not run — external comparisons require a reference to score against.

8 Experimental results

Table 2 maps the full set of quantitative results before the details: every benchmark system, the strongest evaluation available for it, the outcome for **amortix**, the external methods measured on it, and where in the paper the numbers appear. Two conventions hold throughout. Where a reference posterior exists, results are reported as posterior-mean bias (in units of the reference posterior standard deviation) and the width ratio (learned sd over reference sd; 1 is exact, above one conservative, below one overconfident), and the SBC screen is reported as the number of parameters passing the uniformity test. External methods were run only on systems with references, since without a reference the comparison could not be scored.

8.1 Fixed-design gallery

Table 3 summarizes the nine fixed-design systems. One model configuration (the default of Sec. 4) is trained per system on 40,000 simulations for 35 epochs (about 11 minutes per system on one H200 GPU). Calibration is assessed by SBC over 500 synthetic datasets with 200 posterior draws each, and reported in two forms. *Coverage error* asks the practical question: if the posterior’s central $q\%$ interval is trusted, how often does it actually contain the truth? The table reports the mean absolute deviation of empirical from nominal coverage across levels — zero is perfect, and at 500 datasets pure sampling noise alone




Figure 2: Calibration of the gallery models, read the way a practitioner would use them. Left: for each system, the fraction of the 500 test datasets in which the true parameter fell inside the posterior’s central $q\%$ credible interval, plotted against q and averaged over the system’s parameters; a perfectly calibrated posterior follows the diagonal, and the gray band is the spread attributable to having only 500 datasets. All nine systems track the diagonal within the band. Right: the same curves for the individual parameters of the least calibrated system of this run (Cox–Ingersoll–Ross); the b curve shows the ≈ 0.15 sd location shift discussed in the text. Rank-level detail for all 32 parameters is given in Appendix A (Fig. 8).

produces about 1–1.5 percentage points, so the observed mean of 1.4 pp is at the resolution limit of the test. *SBC pass* counts parameters whose rank distribution is consistent with uniformity at $p > 0.05$; with 32 tests, one or two rejections are expected even from a perfectly calibrated model. Of 32 parameters, 31 pass. The failing case in this run is the mean-level parameter b of Cox–Ingersoll–Ross, and the rank data show what the failure is: the truth falls above the posterior median in 56% of datasets instead of 50%, i.e. a location shift of roughly 0.15 posterior sd, while interval coverage remains intact (48% and 88% at nominal 50% and 90%). In an independent earlier training run of the same configuration, 31 of 32 parameters passed as well but the marginal case was a different one (stochastic Lotka–Volterra β , $p = 0.04$, with Cox–Ingersoll–Ross passing at $p \geq 0.33$): defects of 0.1–0.2 sd sit at the resolution limit of SBC at this size, so the identity of the marginal case moves between retrainings — one reason SBC alone is not used as a verdict. On the systems with exact references, the learned posteriors carry 89–90% of the available information (width relative to the exact posterior). Per-parameter p -values for the run shown are listed in Appendix A. Figure 2 shows the calibration directly in the form a user consumes it: empirical against nominal coverage of the central credible intervals, for all nine systems.

8.2 Posterior geometry against validated references

Scalar summaries answer whether the posterior is calibrated on average; Fig. 3 shows the posteriors themselves. For every system with an exact or validated reference, the joint posterior of the trained model is overlaid on the reference for one representative test dataset — the sanity systems (linear-Gaussian, GBM, OU), where a mistake would be inexcusable, and the hard ones: the Merton jump-diffusion, where volatility and jump rate compete to explain the same variance; pharmacokinetics at six irregular blood draws, the data-poor clinical regime; and Fisher–KPP, whose posterior concentrates on the curved ridge $Dr \approx \text{const}$ fixed by the observable front speed. The learned posteriors reproduce location, scale, orientation, and — where present — the curvature of the ridge. This figure, together with the calibration curves of Fig. 2 and the summary of Table 2, is the evidence for the claim that one small default model recovers full posteriors across the suite; the remaining subsections quantify the edges.

8.3 Variable designs

Variable-design zoo. With the class rule and training procedure fixed, six further systems were trained without per-system adjustments. Heston (five parameters, latent volatility): all 25 SBC cases pass, and pair tokens give no advantage over point tokens, consistent with the class boundary (Sec. 4.3). Merton jump-diffusion: 25 of 25 cases after the heavy-tail treatment of Sec. 9, including dense-design cases that

system	d	coverage err.	SBC pass	cov. 50%	cov. 90%
linear-Gaussian	4	1.3 pp	4/4	51%	90%
Ornstein–Uhlenbeck	2	1.8 pp	2/2	47%	88%
SEIR	5	1.5 pp	5/5	49%	90%
geometric Brownian motion	2	1.5 pp	2/2	50%	90%
Cox–Ingersoll–Ross	3	1.5 pp	2/3	49%	88%
double well	3	0.8 pp	3/3	50%	90%
stochastic Lotka–Volterra	4	1.9 pp	4/4	48%	88%
FitzHugh–Nagumo	4	1.1 pp	4/4	51%	90%
SINDy-SDE (cubic drift)	5	1.4 pp	5/5	50%	89%
total / mean	32	1.4 pp	31/32	49%	89%

Table 3: Fixed-design gallery: one default model per system, 40,000 simulations, 35 epochs. Coverage error is the mean absolute deviation of central-interval coverage from its nominal level — zero is perfect, and pure sampling noise at 500 test datasets contributes about 1–1.5 pp, so values in this range indicate calibration at the resolution limit of the test. “SBC pass” counts parameters consistent with uniformity at $p > 0.05$; with 32 tests, one or two rejections are expected by chance even from an exact posterior.

failed at $p < 10^{-3}$ before it; the residual width sits on the budget curve below. Hénon–Heiles (parameters of the potential of Lubich et al. 2015, observing one coordinate at random times): calibrated at every design size down to $K = 6$, where the posterior widens honestly rather than degrading. Hodgkin–Huxley (Hodgkin and Huxley, 1952): calibrated on mixed designs; one dense-design case shows the budget signature. Pharmacokinetics: at $K = 6$ irregular blood draws, bias -0.01 to -0.06 sd and widths 1.09–1.33 against exact-likelihood MCMC. Fisher–KPP: the posterior concentrates on the ridge $D \cdot r \approx \text{const}$; the residual along-ridge bias decays with budget (next subsection).

8.4 Residuals scale as power laws in the training budget

Accuracy against the exact references is a smooth, monotone function of the training budget, and we report it systematically rather than at a single operating point. Figure 4 shows posterior width and bias against the reference for three systems (GBM, OU, Fisher–KPP with the fixed design of Sec. 8.5) over a grid of budgets from 5,000 to 40,000 simulations with optimizer steps scaled proportionally. Posterior widths converge to one from above on all three systems (Fisher–KPP: $1.45 \rightarrow 1.01$). Posterior-mean biases behave differently by identifiability: on GBM and OU they stay within ± 0.1 sd throughout, while on Fisher–KPP the bias lies along the weakly identified ridge direction and does not decay monotonically — independently trained models land at different points within ± 0.5 sd along the ridge, consistent with the flat direction carrying the weakest training signal (Sec. 3). Under the variable-design training procedure, which averages every optimizer step over fresh observation designs, the same ridge bias does decay with budget. Beyond it, the two residual defects of the variable-design suite — inflated width on the hardest Merton case and along-ridge bias on Fisher–KPP under random designs — respond to the same budget lever (Fig. 5). The Merton width excess follows $\propto 1/\text{budget}$ ($0.246 \rightarrow 0.120 \rightarrow 0.059$ at $\times 1, \times 3, \times 6$); the KPP ridge bias follows $\approx \text{budget}^{-1/2}$ ($+0.41 \rightarrow +0.23 \rightarrow +0.14$ sd along the ridge, with the sharply identified front-speed component decaying in parallel). Two negative controls exclude the obvious alternatives: doubling network width and depth at fixed budget changes the Merton case by less than 0.02 in width ratio (1.103 vs 1.120 at $\times 3$), and refining the sampling ODE solver from 20 midpoint steps to 60 RK4 steps changes nothing measurable. The weakly identified posterior direction carries the weakest regression signal in (3) and converges last; we find no evidence of a floor within the measured range.

8.5 Comparison with existing implementations

Protocol. We compare against the directly comparable amortized estimators available in maintained packages: NPE with its default masked autoregressive flow (Papamakarios et al., 2017) and FMPE (Wildberger et al., 2023) from `sbi` 0.27 (Boelts et al., 2024); the coupling-flow and flow-matching networks of BayesFlow 2.0 (Radev et al., 2020); and, because it is the closest architecture to ours in the literature, a port of the Simformer (Gloeckler et al., 2024) built from its authors’ minimal example (transformer–diffusion over the joint, random condition masks, their architecture and SDE settings, full 75,000-step training). All arms receive the same simulation budget (20,000 draws from the same prior and simulator), run with their package defaults and their own stopping rules, and are compared on exactly three axes:




Figure 3: Joint posteriors of **amortix** (filled contours: 39% and 86% highest-density regions) against exact or validated-MCMC reference posteriors (open contours, same levels) on one test dataset per system; crosses mark the true parameters. Top row: systems with closed-form or cheap exact references. Bottom row: systems where the reference itself required exact-likelihood MCMC with two-chain validation — including the ridge-shaped Fisher–KPP posterior, where the model reproduces the curvature, not just the moments.

training cost (wall-clock time to fit at this budget), accuracy (posterior-mean bias and width ratio per parameter against the reference posterior, as in Sec. 6.2), and inference cost (wall-clock time per dataset for 2,000 posterior samples). Because these estimators consume fixed-length vectors, the comparison uses fixed observation designs. Three systems are used, chosen to separate the possible explanations of any performance gap: a multiplicative process where input conditioning is hostile (GBM, exact conjugate reference), an additive process where it is comparatively benign (OU, exact-likelihood MCMC reference, two validated chains per dataset), and a PDE where every likelihood evaluation requires a forward solve (Fisher–KPP), so that the cost relationship between reference and amortized inference is inverted. Where informative, external arms are additionally given hand-derived inputs — log-prices, or per-gap log-returns, the exact coordinates of the GBM likelihood factorization — to separate the effect of the input representation from the generative engine.

Multiplicative process (GBM). Table 4 compares all methods at their defaults on raw price data; Table 5 repeats the external methods with hand-derived input coordinates; Fig. 6 shows the accuracy axis of both tables at a glance. On raw data every external configuration inflates the volatility posterior — widths 1.85 to 3.16 — irrespective of the generative engine; **sbi** itself emits a warning about precision loss during z-scoring on this data, which is a fair description of the mechanism (Sec. 9). Hand-derived coordinates do not produce convergent behavior (Table 5): with identical log-return inputs, NPE remains 44% over-dispersed, FMPE is nearly exact (1.09), both BayesFlow arms become overconfident (widths 0.42–0.85), and the Simformer port acquires a -0.48 sd bias. Each of these defects is invisible without an exact reference, and each lies within or near the blind region of SBC at practical sizes (Sec. 6). At fixed input coordinates the flow-matching engine outperforms the autoregressive flow (1.09 vs 1.44), consistent with the engine choice of Sec. 3, but no external arm matches the learnable representation operating on raw data (1.05). The verdicts are stable across devices and seeds: re-running the **sbi** and BayesFlow arms on an H200 GPU with fresh seeds reproduces every failure direction (NPE 2.48, FMPE 2.30, BayesFlow coupling 1.80 on raw input; BayesFlow overconfident at 0.40–0.72 on log-returns).




Figure 4: Convergence to the reference posterior with training budget on three systems (columns): posterior width ratio (top) and posterior-mean bias (bottom) per parameter, over budgets of 5,000–40,000 simulations with optimizer steps scaled proportionally. References: exact conjugate posterior (GBM), exact-likelihood MCMC (OU), PDE-likelihood MCMC (Fisher-KPP, fixed $K = 40$ design).

Additive control (Ornstein–Uhlenbeck). The same protocol on OU — an additive process with no floating scale — separates the candidate explanations of the GBM results (Table 6). The reference is exact-likelihood MCMC (per-gap Euler–Maruyama transitions plus the stationary initial density; 62 ms per chain; two-chain validation: worst mean discrepancy 0.29 sd, mean 0.08 sd over 100 datasets). The external estimators improve markedly relative to GBM: the volatility bias is no longer significant ($+0.19 \pm 0.19$ for NPE, $+0.08 \pm 0.12$ for FMPE), and the width inflation drops from 2.5–3.2 to 1.9 (NPE) and 1.6 (FMPE). This is the pattern the conditioning analysis predicts: the form-1 pathology of Sec. 9 is absent by construction, and what remains is the general difficulty of assembling quadratic-variation statistics from a raw high-dimensional vector — present but smaller. `amortix` reads 1.03/1.05 on the same data.

Simulator-limited regime (Fisher–KPP). The third system inverts the cost structure of the first two: every likelihood evaluation requires solving the forward model. A one-dimensional reaction–diffusion equation is inexpensive in absolute terms (16 ms per solve here); what makes the case informative is the *ratio* between reference and amortized inference, because both sides scale linearly with the solve cost — a chain of n likelihood evaluations costs n solves per dataset whatever the forward model, while the amortized cost per dataset does not involve the simulator at all. The measured ratio below transfers to genuinely expensive simulators, where the same reference sampler would take days per dataset. A single fixed design of $K = 40$ random (time, sensor) points is shared by all datasets and methods; the reference is the PDE-likelihood sampler of Sec. 6.2 with two validated chains per dataset (271 s per dataset for the pair; two-chain validation over 32 datasets: worst mean discrepancy 0.22 sd, mean 0.07 sd). Table 7 gives the results. `amortix` matches the reference (1.02/1.04 width, biases within 0.13 sd) at 31 ms per dataset; its entire training cost (379 s) equals the reference cost of fewer than three datasets, so amortization pays for itself almost immediately in any repeated-inference setting. NPE at defaults shows moderate width inflation (1.26–1.28); FMPE at defaults fails outright on the diffusivity ($+7.3 \pm 2.7$ sd bias at $12.5\times$ width) while remaining reasonable on the reaction rate — a further instance of default configurations behaving discontinuously across parameters, and detectable here only because the expensive reference was computed.

Scope of the comparison. Every external arm ran at its package defaults under our simulation budget; specialists could improve each with problem-specific embedding networks or summary statistics.




Figure 5: The two residual defects of the suite against training budget (simulations and steps scaled together; model size fixed). Left: Merton width excess on the hardest evaluation case, with the $1/\text{budget}$ guide. Right: Fisher-KPP posterior-mean biases in the ridge coordinate system, with the $\text{budget}^{-1/2}$ guide.

method	train	inference	μ : bias / width	σ : bias / width
amortix CFM	488 s	426 ms	+0.00 / 1.00	-0.00 / 1.05
sbi NPE	29 s	13 ms	-0.02 / 1.06	+0.45 $\pm$ 0.19 / 2.64
sbi FMPE	27 s	370 ms	-0.03 / 1.03	+0.20 $\pm$ 0.16 / 2.46
BayesFlow 2 coupling	384 s	4 ms	-0.04 / 0.95	-0.13 $\pm$ 0.11 / 1.85
BayesFlow 2 flow matching	220 s	2220 ms	+0.06 / 1.00	-0.36 $\pm$ 0.11 / 1.99
Simformer port*	591 s	1057 ms	-0.04 / 1.08	-0.30 $\pm$ 0.20 / 3.16
exact-likelihood MCMC	—	42 ms	reference (chain check: -0.04 / +0.01)	

Table 4: GBM, all methods on raw price data at package defaults: training cost, accuracy against the exact posterior (bias in reference-sd units; width ratio, 1 exact), and inference cost per dataset for 2,000 samples. Fixed 95-point design, 20,000-simulation budget, 200 test datasets; timings on one laptop CPU. *Simformer is research code ported from its authors’ minimal example and run on an H200 GPU: its timing entries are not device-comparable with the rest, its accuracy entries are.

The comparison is designed to measure exactly that dependence: what accuracy is delivered when the input representation is not hand-engineered per problem. The Simformer arm in particular is a port of a minimal example rather than a tuned configuration, and its numbers should be read as representative of that recipe, not of the method’s ceiling. Conversely, the variable-design capability of Sec. 8.3 is not exercised here at all, since the external estimators consume fixed-length inputs by default.

8.6 Ablations: what each architecture component contributes

Table 8 and Fig. 7 compare encoder variants for design amortization on GBM, trained at $K \sim \log\text{-uniform}[2, 128]$ random-time designs and evaluated against the exact per-design posterior (96 datasets per bucket). Three observations organize the table. First, index-based positions combined with in-place subsampling during training (leaving kept tokens at spread positions while inference packs contiguously) produce biases of +1.1 to +3.3 posterior sd with no individually incorrect component — the positional pathology of Sec. 4.2; repacking or replacing positions with time phases removes the bias entirely. Second, all point-token variants share inflated widths at dense designs, in proportion to how poorly their pooled representation encodes $\log K$ (rightmost column; the probe of Sec. 4.3) — including two attention mechanisms specifically constructed to exploit continuous time (a learned per-head time-kernel logit bias in the spirit of relative-position methods, and edge-message attention carrying gap-modulated increments on all pairs). Third, pair tokens with time-phase attention, plus the training procedure of Sec. 5, bring every bucket to within 7% of the exact width.

8.7 Computational cost

The default model trains in minutes to tens of minutes per system: on one H200 GPU, a variable-design zoo system at base budget (20,000 simulations, 12,000 steps) takes 8–12 minutes, and the $\times 6$ budget behind the sharpest numbers of Fig. 5 takes about 50 minutes; on a laptop CPU the fixed-design gallery

method	input given	μ : bias / width	σ : bias / width
sbi NPE	log-prices	$-0.04 / 1.01$	$+0.32 \pm 0.17 / 2.29$
sbi NPE	log-returns	$+0.07 / 1.07$	$+0.13 \pm 0.10 / 1.44$
sbi FMPE	log-returns	$+0.01 / 1.03$	$-0.02 \pm 0.07 / 1.09$
BayesFlow 2 coupling	log-returns	$-0.01 / 0.42$	$-0.15 \pm 0.05 / 0.54$
BayesFlow 2 flow matching	log-returns	$-0.04 / 0.46$	$-0.30 \pm 0.06 / 0.85$
Simformer port	log-returns	$-0.08 / 1.10$	$-0.48 \pm 0.02 / 1.36$

Table 5: Input sensitivity of the external methods on GBM: the same tools as Table 4 given hand-derived coordinates — log-prices, or the per-gap log-returns that constitute the exact coordinates of the likelihood factorization. Widths below one are overconfident. **amortix** appears only in Table 4: it receives raw data by construction.




Figure 6: The accuracy axis of Tables 4 and 5 at a glance: volatility posterior width as a ratio to the exact posterior, so 1 is exact, larger is too wide, and values below one are overconfident. Annotations give the absolute posterior-mean bias in units of the exact posterior standard deviation.

systems take 8–45 minutes. Inference draws 2,000 posterior samples for a dataset in 426 ms on the laptop CPU and 72 ms (batched; 135 ms single) on the GPU. GPU and CPU runs produce bit-identical posteriors because all stochastic generation is kept on a CPU random number generator by design, so a device change cannot silently alter results. For orientation against non-amortized inference: on GBM, whose likelihood is a three-line formula, adaptive Metropolis costs 42 ms per dataset and amortization buys nothing for a single dataset; on Fisher–KPP, where one likelihood evaluation is one PDE solve (16 ms), the same reference sampler costs minutes per dataset (measured in Table 7); and for Heston or Hodgkin–Huxley no likelihood formula is available at any price. The regime where amortization pays is determined by the likelihood cost, the number of datasets, and — as Table 4 indicates — by whether a trustworthy input representation for a conditional density estimator is available without one.

9 Failure analysis: conditioning of the encoder input

Every accuracy failure encountered during the development of this package — including, we argue, the external failures of Table 4 — reduces to a single cause appearing in three forms: the conditioning of the numbers presented to the set encoder. We document the three forms with their measurements and repairs; all three repairs are learnable modules rather than problem-specific transformations.

Form 1: floating multiplicative scale. On GBM, raw price tokens force the network to divide squared increments by squared levels internally; with global standardization this manifests as a volatility bias of $+0.48$ posterior sd (confirmed against the closed-form posterior) that no amount of training removes. The repair is the monotone mixture reparametrization of Sec. 4.3, which learns the log coordinate; the Yeo–Johnson ablation there shows why the parametrization, not the hypothesis class, is the operative choice. External estimators exhibit the same form on the same data (Table 4, raw-price rows).

method (raw values as input)	train	inference	θ : bias / width	σ : bias / width
amortix CFM	183 s	32 ms	−0.01 / 1.03	+0.01 / 1.05
sbi NPE	128 s	10 ms	+0.07 / 1.15	+0.19 ± 0.19 / 1.88
sbi FMPE	87 s	147 ms	+0.11 / 1.01	+0.08 ± 0.12 / 1.60
exact-likelihood MCMC	—	124 ms	reference (two chains per dataset)	

Table 6: Ornstein–Uhlenbeck, fixed 74-point design, 20,000-simulation budget, 100 test datasets, one H200 GPU (MCMC on its CPU; time given for the two validation chains). The additive control: external biases vanish and widths improve but remain inflated; compare Table 4.

method (40 observed values as input)	train	inference	D : bias / width	r : bias / width
amortix CFM	379 s	31 ms	−0.05 / 1.02	+0.13 / 1.04
sbi NPE	218 s	7 ms	+0.19 ± 0.12 / 1.28	−0.06 / 1.26
sbi FMPE	115 s	405 ms	+7.3 ± 2.7 / 12.5	−0.25 ± 0.10 / 1.19
PDE-likelihood MCMC	—	271 s	reference (two chains per dataset)	

Table 7: Fisher–KPP, fixed $K = 40$ spatio-temporal design, 20,000-simulation budget, 32 test datasets, one H200 GPU (MCMC on its CPU pool). One likelihood evaluation is one PDE solve, so the reference costs 271 s per dataset; the training cost of the amortized model equals the reference cost of fewer than three datasets.

Form 2: cardinality blindness. Normalized attention erases set size, which controls posterior width under design amortization; the probe measurements and the $\log K$ /register-token repairs are given in Sec. 4.3, with the arm-by-arm consequences in Table 8.

Form 3: heavy tails. On the Merton jump-diffusion, squared increments have sample standard deviation 3.1×10^4 against a diffusion signal of order 10^{-2} –1; running feature normalization never converges, and the flow trains against a drifting representation, producing a +1.12 sd volatility bias. Encoder probing (Sec. 6.3) shows the information is extracted — the failing configuration encodes robust volatility statistics better than the healthy one — so the repair targets representation stability: logarithmic compression of increment features (near-identity on clean diffusions since $\log(1+x) \approx x$ for small x) restores all 25 evaluation cases. A precision follow-up against the Poisson-mixture reference then isolated a smaller (+0.27 sd) residual whose mechanism is a jump-classification threshold that floats with the unknown volatility; a per-dataset set-conditioning module (pooled statistics modulating feature channels, initialized to the identity) removes it as effectively as a hand-crafted per-dataset rescaling ($+0.127 \pm 0.055$ vs $+0.094 \pm 0.054$ sd, statistically indistinguishable). This was one of three such paired comparisons in the project (log coordinate, set size, robust scaling); in all three the learnable module matched the hand-crafted repair once initialized at the working configuration. We summarize the pattern as a design principle: the hand-crafted fix identifies the mechanism and the initialization; the learnable module placed there is what ships.

Three implementation conditions proved necessary for the set-conditioning module and are recorded here because each was found through a failure: the modulation must be zero-initialized (training must start from the working baseline); it must act after feature normalization, not before (the reverse ordering re-creates the representation drift it is meant to cure); and pooled set statistics must exclude padding positions (a masking error here made padded and packed representations of the same dataset inequivalent, detected only because SBC and the exact-reference probe disagreed).

10 Discussion and limitations

The evidence assembled here supports a specific claim: for dynamical-systems inverse problems, the input representation of an amortized posterior estimator is not a preprocessing detail but the dominant accuracy factor, and it can be made part of the model rather than part of the problem specification. The comparison of Sec. 8.5 shows the cost of the alternative — fixed representations with global standardization — across four generative engines, on a problem where the correct answer is known in closed form. The OU control separates the mechanisms: on additive data the bias of the external estimators disappears and their width inflation drops by roughly a factor of two but does not vanish, so the conditioning pathology accounts for the dominant part of the failure and a smaller representation-learning gap accounts for the rest.

encoder variant	σ width ratio at K				$R^2(\log K)$
	5	9	20	100	
point tokens, index positions, spread subsampling	0.94*	1.13*	1.42*	2.08*	—
point tokens, index positions, contiguous	1.05	1.17	1.13	1.73	0.957
point tokens, time-kernel attention bias	1.03	1.14	1.31	2.51	0.712
point tokens, edge-message attention	1.03	1.15	1.29	2.45	0.498
point tokens, time phases	1.01	1.10	1.14	1.95	—
+ explicit $\log K$ feature	1.01	1.05	1.08	1.49	—
pair tokens, time phases	1.02	1.02	1.02	1.27	0.985
pair tokens + full training procedure	1.06	1.06	1.05	1.07	—

Table 8: GBM with variable designs: posterior width for the volatility σ against the exact per-design posterior (96 datasets per design-size bucket) and, where measured, the probe R^2 for $\log K$ from the pooled encoder output. The table tracks σ because it is the parameter whose posterior width is governed by the observation density — the drift μ is weakly informed at every K , and all variants recover its posterior within a few percent, so it does not discriminate between architectures. *This row also carries biases of +1.09 to +3.30 posterior sd (the positional train/test mismatch of Sec. 4.2); all other rows have biases consistent with zero.




Figure 7: The data of Table 8 in graphical form: excess posterior width for the GBM volatility — the width ratio against the exact posterior minus one, so 0.1 means a posterior 10% too wide and lower is better — as a function of the number of observation points K (both axes logarithmic). Dotted curves: point-token variants with no set-size information, which degrade as designs become dense; dashed: partial repairs; solid: increment tokens, and increment tokens combined with the training procedure of Sec. 5, which is flat in K .

Limitations. (i) The remaining inaccuracies of the suite decay as measured power laws in training compute (Fig. 5); we have not searched for methods that improve the exponents. (ii) The power limits of SBC imply that small residuals can only be certified where an exact or validated reference exists; systems without tractable likelihoods inherit the weaker screen, and carrying reference-equipped systems alongside them in one suite is our current mitigation. (iii) The comparison of Sec. 8.5 uses package defaults; it measures the out-of-the-box regime, not the ceiling of any method. (iv) Bayesian experimental design — the intended consumer of design amortization — is not yet demonstrated; the training procedure of Sec. 5 was built so that the required object $p(m | S)$ is already coherent across subsets S . (v) No real-data study with posterior predictive checks is included yet. (vi) All priors are uniform on boxes; the probit normalization (2) extends to any factorized prior with tractable CDFs, but correlated priors would require a different normalization.

A Benchmark system definitions

All SDEs are integrated by Euler–Maruyama on the stated fine grid; ODEs by RK4 (Hodgkin–Huxley by a semi-implicit scheme for the gating variables); the PDE by finite differences. Fixed-design observations are token sets over the union of the stated channels; “fast/slow” channels observe every step over a

short window and every j -th step over the horizon respectively. Priors are uniform on the stated boxes. Classical estimators (`sota()`) are listed per system.

Linear-Gaussian. $y = Am + \varepsilon$, $\varepsilon \sim \mathcal{N}(0, 0.5^2 I_6)$, with a fixed correlated design matrix $A \in \mathbb{R}^{6 \times 4}$; $m \in [-3, 3]^4$. Exact Gaussian posterior; the classical estimator is the posterior mean itself. This system checks the flow engine in isolation from the encoder.

Ornstein–Uhlenbeck. $dX = -\theta X dt + \sigma dW$, stationary start $X_0 \sim \mathcal{N}(0, \sigma^2/2\theta)$; $dt = 0.02$, 500 steps; channels: fast (1×50), slow (20×24). Prior $\theta \in [0.3, 3]$, $\sigma \in [0.2, 1.5]$. Classical: exact MLE from per-gap Gaussian transitions.

Geometric Brownian motion. $dS = \mu S dt + \sigma S dW$, $S_0 = 1$; $dt = 0.01$, 500 steps (five “years”); channels: fast (1×60), slow (10×40). Prior $\mu \in [-0.2, 0.4]$, $\sigma \in [0.1, 0.6]$. Classical: exact MLE on log-returns. Reference: conjugate posterior (Appendix B).

Cox–Ingersoll–Ross. $dX = a(b - X) dt + \sigma \sqrt{X} dW$; $dt = 0.02$, 500 steps; channels as OU. Prior $a \in [0.3, 3]$, $b \in [0.3, 1.5]$, $\sigma \in [0.1, 0.5]$. Classical: Euler pseudo-likelihood MLE (Cox et al., 1985).

Double well. $dX = (\theta_1 X - \theta_2 X^3) dt + \sigma dW$; $dt = 0.01$, 1000 steps, noise strong enough for barrier crossings; channels: fast (1×80), slow (25×36). Prior $\theta_{1,2} \in [0.5, 3]$, $\sigma \in [0.4, 1.2]$. Classical: Kramers–Moyal regression.

Stochastic Lotka–Volterra. Prey/predator pair with multiplicative noise, $dx_1 = (\alpha x_1 - \beta x_1 x_2) dt + 0.1 x_1 dW_1$, $dx_2 = (\delta x_1 x_2 - \gamma x_2) dt + 0.1 x_2 dW_2$; both components observed; $dt = 0.01$, 600 steps; channels: fast (1×50), slow (15×40). Prior $\alpha, \gamma \in [0.8, 1.5]$, $\beta, \delta \in [0.4, 1.2]$. Classical: multi-start nonlinear least squares.

FitzHugh–Nagumo. $\dot{v} = v - v^3/3 - w + I + \text{noise}$, $\dot{w} = \epsilon(v + a - bw)$; only v observed at 25 evenly spaced times over horizon 40 ($dt = 0.05$), observation noise 0.05. Prior $a, b \in [0.5, 0.9]$, $\epsilon \in [0.05, 0.3]$, $I \in [0, 0.5]$. Classical: multi-start nonlinear least squares on the deterministic solve.

SEIR with mortality. Five-compartment stochastic epidemic (two-phase contact rate $\beta_1 \rightarrow \beta_2$, incubation α , recovery γ_r , mortality γ_d); infected and deceased compartments observed at 20 days over a 60-day horizon ($dt = 0.1$), observation noise 0.004. Prior in Table 9. Classical: nonlinear least squares on (I, D) .

SINDy-SDE. $dX = (c_0 + c_1 X + c_2 X^2 + c_3 X^3) dt + \sigma dW$ (in the spirit of Brunton et al., 2016); $dt = 0.01$, 1000 steps; channels: fast (1×80), slow (20×45). Prior $c_0, c_2 \in [-0.5, 0.5]$, $c_1 \in [-2, 0]$, $c_3 \in [-1, -0.1]$, $\sigma \in [0.3, 0.8]$. Classical: SINDy-style least squares on Kramers–Moyal estimates.

Heston (variable designs). $dS = \mu S dt + \sqrt{v} S dW_1$, $dv = \kappa(\bar{v} - v) dt + \xi \sqrt{v} dW_2$, $\text{corr}(dW_1, dW_2) = \rho$ (Heston, 1993); price observed, volatility latent; $dt = 0.01$, 500 steps, designs of $K \in [2, 128]$ random times. Prior $\mu \in [-0.1, 0.3]$, $\kappa \in [0.5, 5]$, $\bar{v} \in [0.02, 0.3]$, $\xi \in [0.1, 1]$, $\rho \in [-0.9, 0]$.

Merton jump-diffusion (variable designs). $dS/S = \mu dt + \sigma dW + d(\sum_i (e^{J_i} - 1))$ with Poisson jump rate λ and $J_i \sim \mathcal{N}(\mu_J, s_J^2)$ (Merton, 1976); $dt = 0.01$, 500 steps, $K \in [2, 128]$. Prior $\mu \in [-0.1, 0.3]$, $\sigma \in [0.1, 0.5]$, $\lambda \in [0.2, 3]$, $\mu_J \in [-0.3, 0.3]$, $s_J \in [0.05, 0.4]$. Reference: Poisson-mixture transition density (Appendix B).

Hénon–Heiles (variable designs). Two-dimensional Hamiltonian oscillator with the coupling potential used in Lubich et al. (2015), frequencies (ω_1, ω_2) and coupling λ ; q_1 observed at random times with noise; $dt = 0.05$, 800 steps, $K \in [4, 64]$. Prior $\omega_{1,2} \in [0.7, 1.3]$, $\lambda \in [0.02, 0.22]$.

Hodgkin–Huxley (variable designs). Standard four-dimensional membrane model (Hodgkin and Huxley, 1952); voltage observed at random times; parameters $(g_{\text{Na}}, g_K, g_L, I)$ with prior $[60, 180] \times [10, 50] \times [0.05, 0.5] \times [4, 12]$; 3000 integration steps, $K \in [4, 96]$.

system	per-parameter SBC p -values					
linear-Gaussian	m_1 : 0.06	m_2 : 0.60	m_3 : 0.38	m_4 : 0.56		
Ornstein-Uhlenbeck	θ : 0.28	σ : 0.77				
SEIR	β_1 : 0.33	β_2 : 0.22	α : 0.61	γ_r : 0.34	γ_d : 0.53	
geometric Brownian motion	μ : 0.26	σ : 0.44				
Cox-Ingersoll-Ross	a : 0.65	b : 0.001	σ : 0.73			
double well	θ_1 : 0.78	θ_2 : 0.54	σ : 0.88			
stochastic Lotka-Volterra	α : 0.07	β : 0.69	δ : 0.08	γ : 0.35		
FitzHugh-Nagumo	a : 0.19	b : 0.15	ϵ : 0.90	I : 0.63		
SINDy-SDE	c_0 : 0.06	c_1 : 0.96	c_2 : 0.46	c_3 : 0.80	σ : 0.50	

Table 9: Per-parameter SBC uniformity p -values for the fixed-design gallery run of Table 3 (500 datasets $\times$ 200 draws; $p > 0.05$ passes). The single failure is marked in bold; see Sec. 8.1 for its behavior across retrains.

Pharmacokinetics (variable designs). One-compartment oral dosing (Bateman function), concentration $c(t) = \frac{D_0 k_a}{V(k_a - k_e)}(e^{-k_e t} - e^{-k_a t})$ with dose $D_0 = 500$; log-normal assay noise (10%); horizon 24 h, $K \in [3, 64]$ irregular draws. Prior $k_a \in [0.5, 4]$, $k_e \in [0.05, 0.4]$, $V \in [20, 100]$. Reference: log-normal likelihood MCMC.

Fisher-KPP (variable designs). $\partial_t \theta = D \partial_{xx} \theta + r \theta(1 - \theta)$ on $[0, 1]$ with reflecting boundaries, 64-point finite-difference grid, Gaussian initial bump; three fixed sensors at $x = 16/63, 32/63, 48/63$ observed at random (time, sensor) pairs with noise 0.02; $dt = 0.01$, 200 steps, $K \in [4, 96]$. Prior $D \in [0.001, 0.01]$, $r \in [2, 10]$. Reference: Gaussian observation model around the deterministic solve; one likelihood evaluation is one PDE solve.

B Reference posteriors

GBM conjugate posterior. With observations at indices $i_0 = 0 < i_1 < \dots$ of the fine grid, log-returns $r_k = \log(S_{i_{k+1}}/S_{i_k})$ over gaps τ_k are independent $\mathcal{N}((\mu - \sigma^2/2)\tau_k, \sigma^2 \tau_k)$. In the coordinates $(b, \sigma^2) = (\mu - \sigma^2/2, \sigma^2)$ the model is a Gaussian location-scale family with sufficient statistics $T = \sum_k \tau_k$, $R_1 = \sum_k r_k$, and $SS_w = \sum_k r_k^2/\tau_k - R_1^2/T$; draws from the standard conjugate posterior ($\sigma^2 = SS_w/\chi_{n-1}^2$, $b \mid \sigma^2 \sim \mathcal{N}(R_1/T, \sigma^2/T)$) are transformed to (μ, σ) and importance-resampled to the uniform box prior with the appropriate Jacobian weight. The sampler was validated against exact-likelihood MCMC on the same designs: over 24 comparisons, worst posterior-mean discrepancy 0.099 sd, width ratios 0.91–1.03.

Adaptive Metropolis. Random-walk Metropolis on the box, proposal $x' = x + s C \xi$, $\xi \sim \mathcal{N}(0, I)$. During burn-in (default $\max(500, n/2)$ iterations) the scalar s is adapted toward a 0.25 acceptance rate in blocks of 50 with a decaying Robbins-Monro gain, and once sufficient burn-in states are collected the shape C is replaced by the Cholesky factor of the empirical covariance rescaled by $2.38/\sqrt{d}$ (Haario et al., 2001); adaptation is frozen for the retained phase. Out-of-box proposals are rejected, which samples the box-truncated target exactly since the proposal is symmetric. Likelihoods: per-gap Euler-Maruyama Gaussian transitions for OU (matching the generative process, plus the stationary density of X_0); the exact GBM log-return density; the Poisson-mixture transition density for Merton (truncated at the numerically negligible jump count); log-normal residuals around the Bateman curve; a Gaussian observation model around the deterministic PDE solve for Fisher-KPP. Every chain used as a reference is paired with a second chain from a well-separated start; between-chain posterior-mean discrepancies are reported alongside the results they certify (Sec. 8.5).




Figure 8: Calibration check for the fixed-design gallery. Procedure: for each system, 500 synthetic datasets are simulated with known parameters; for each dataset the model produces 200 posterior draws, and the rank of the true parameter among the draws is recorded. If the posterior is calibrated, these ranks are uniform. Each curve shows, for one parameter, the deviation of the empirical CDF of normalized ranks from the uniform diagonal; the gray band is the 99% range attributable to sampling noise at 500 datasets. A curve bulging below zero on the left indicates a posterior shifted low relative to the truth, and vice versa. All curves stay within the band except Cox–Ingersoll–Ross b , whose deviation corresponds to the ≈ 0.15 sd location shift discussed in the text.

C Correspondence with package options

in this report	in the package
point tokens	<code>embed="wpoint"</code>
pair tokens + set conditioning	<code>embed="wfilm"</code>
warped values/increments (fixed designs)	<code>embed="wbasis" / "wdiff"</code>
attention phases from physical time	<code>rope="time"</code>
fresh designs each optimizer step	<code>fit(retokenize=...) via make_retokenizer()</code>
variable-design problem wrapper	<code>amortix.designs.DesignProblem</code>
benchmark systems	<code>amortix.problems (DESIGN_Z00)</code>
calibration screen / reference machinery	<code>sbc_design / amortix.mcmc + likelihood factories</code>
comparison scripts (Sec. 8.5)	<code>examples/baseline_{npe,ou,kpp,bayesflow,simformer}.py</code>

All defaults named in this report are selected automatically by `FlowPosterior(problem)`; the non-default options exist as reproducible controls for the comparisons above.

References

- J. Boelts, M. Deistler, M. Gloeckler, et al. sbi reloaded: a toolkit for simulation-based inference workflows. *arXiv:2411.17337*, 2024.
- S. L. Brunton, J. L. Proctor, J. N. Kutz. Discovering governing equations from data by sparse identification of nonlinear dynamical systems. *PNAS* 113(15), 2016.
- J. C. Cox, J. E. Ingersoll, S. A. Ross. A theory of the term structure of interest rates. *Econometrica* 53(2), 1985.
- K. Cranmer, J. Brehmer, G. Louppe. The frontier of simulation-based inference. *PNAS* 117(48), 2020.
- T. Darcet, M. Oquab, J. Mairal, P. Bojanowski. Vision transformers need registers. *ICLR*, 2024.
- M. Gloeckler, M. Deistler, C. Weilbach, F. Wood, J. H. Macke. All-in-one simulation-based inference. *ICML*, 2024.
- D. Greenberg, M. Nonnenmacher, J. Macke. Automatic posterior transformation for likelihood-free inference. *ICML*, 2019.
- H. Haario, E. Saksman, J. Tamminen. An adaptive Metropolis algorithm. *Bernoulli* 7(2), 2001.
- J. Hermans, A. Delaunoy, F. Rozet, A. Wehenkel, V. Begy, G. Louppe. A trust crisis in simulation-based inference? Your posterior approximations can be unfaithful. *TMLR*, 2022.
- S. L. Heston. A closed-form solution for options with stochastic volatility with applications to bond and currency options. *Rev. Financial Studies* 6(2), 1993.
- A. L. Hodgkin, A. F. Huxley. A quantitative description of membrane current and its application to conduction and excitation in nerve. *J. Physiology* 117(4), 1952.
- M. Horn, M. Moor, C. Bock, B. Rieck, K. Borgwardt. Set functions for time series. *ICML*, 2020.
- Y. Lipman, R. T. Q. Chen, H. Ben-Hamu, M. Nickel, M. Le. Flow matching for generative modeling. *ICLR*, 2023.
- C. Lubich, I. V. Oseledets, B. Vandereycken. Time integration of tensor trains. *SIAM J. Numer. Anal.* 53(2), 2015.
- J.-M. Lueckmann, J. Boelts, D. Greenberg, P. Gonçalves, J. Macke. Benchmarking simulation-based inference. *AISTATS*, 2021.
- R. C. Merton. Option pricing when underlying stock returns are discontinuous. *J. Financial Economics* 3(1-2), 1976.
- G. Papamakarios, I. Murray. Fast ϵ -free inference of simulation models with Bayesian conditional density estimation. *NeurIPS*, 2016.

- G. Papamakarios, T. Pavlakou, I. Murray. Masked autoregressive flow for density estimation. *NeurIPS*, 2017.
- E. Perez, F. Strub, H. de Vries, V. Dumoulin, A. Courville. FiLM: Visual reasoning with a general conditioning layer. *AAAI*, 2018.
- W. Peebles, S. Xie. Scalable diffusion models with transformers. *ICCV*, 2023.
- S. T. Radev, U. K. Mertens, A. Voss, L. Ardizzone, U. Köthe. BayesFlow: Learning complex stochastic models with invertible neural networks. *IEEE Trans. Neural Netw. Learn. Syst.* 33(4), 2020.
- S. N. Shukla, B. M. Marlin. Multi-time attention networks for irregularly sampled time series. *ICLR*, 2021.
- J. Su, M. Ahmed, Y. Lu, S. Pan, W. Bo, Y. Liu. RoFormer: Enhanced transformer with rotary position embedding. *Neurocomputing* 568, 2024.
- S. Talts, M. Betancourt, D. Simpson, A. Vehtari, A. Gelman. Validating Bayesian inference algorithms with simulation-based calibration. *arXiv:1804.06788*, 2018.
- J. Wildberger, M. Dax, S. Buchholz, S. R. Green, J. H. Macke, B. Schölkopf. Flow matching for scalable simulation-based inference. *NeurIPS*, 2023.
- I.-K. Yeo, R. A. Johnson. A new family of power transformations to improve normality or symmetry. *Biometrika* 87(4), 2000.